

Stitcher Capacity Analysis:  
Charter Internal Note.  
Draft Version 0.0.0

Brad Schoenrock  
Video Operations Engineering  
Charter Communications  
Greenwood Village, CO

July. 2019

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>QAM Bandwidth Availability &amp; Analysis</b>	<b>3</b>
<b>3</b>	<b>Conclusion</b>	<b>4</b>
<b>A</b>	<b>Appendix A Example QRM Configuration</b>	<b>6</b>
<b>B</b>	<b>Appendix B Customers by SGID</b>	<b>7</b>

## 1 Introduction

Stuff and things...

Spec-Guide QAM bandwidth was initially a concern brought up as a possible cause of Guide Unavailable, Failure to Tune messages, and a potentially call volume generating problem. Spec-Guide QAM bandwidth is a limited resource, which scales with the number of customers, so the capacity of the Spec-Guide QAMs were also a concern to be investigated. The architecture of the platform dictates that each service group gets exactly two Spec-Guide QAMs to serve it, with no more or less. This creates a situation where the RF bandwidth of the QAMs dictate the maximum number of customers a service group can support. Since the QAMs themselves can't output more bandwidth (They are fixed devices that don't lend themselves to upgrades) and our architecture dictates that we can't add more than two QAMs per service group, this means that the only lever that is available to resolve such issues is managing the size and number of our service groups. We (ISP) does this routinely with service group splits and merges. When service groups get too large they are split into two different service groups, and when small service groups are near each other they are collapsed into one service group.

On a service group by service group level the QAM bandwidth is managed on the CSM by the QAM Resource Manager (QRM). The QRM is set up through configuration files, topology.xml and programs.xml. These config files define the total bandwidth output capability of the QAM for the QRM, the output frequency of the QAM, and the bandwidth expected for each type of stream (SD vs. HD, etc...) in both steady state operations as well as a maximum bandwidth usage. In this way we can monitor estimated QAM usage and ensure there aren't any overruns. As overruns would be approached, the QRM uses smartmuxing in order to serve extra sessions when current sessions are not using up their allotted bandwidth.

## 2 QAM Bandwidth Availability & Analysis

Stuff and things...

Each QAM has 38.8MBps of bandwidth, 37.5MBps is usable for delivery of Spec-Guide sessions. With two QAMs per service group this means each service group has 75MBps of bandwidth for delivery of Spec-Guide sessions.

The QRM is configured based on actual session usage, a sample configuration can be seen in Appendix /refAPPENDIX-QRMConfig. The highest definition mpeg streams allocate 2.5MBps steady state, and a peak of 6MBps of RF bandwidth. By monitoring the throughput of the APEX by eye in prime time these settings are matching usage. Average session bandwidth is between 2-2.5MBps which is below 2.5 since not every session is HD.

In order to define a functional capacity we need to ensure enough bandwidth to serve another session. For this reason 6MBps of bandwidth should be reserved on each QAM. The remaining 63MBps of QAM bandwidth on the service group

can then serve 26 HD sessions at the configured 2.5MBps each. Maximum concurrency rate of our customer base was measured to be 4% in stitcher capacity analysis, and while that would fluctuate more on a service group basis, it still serves as a baseline for system usage. This leads to a capacity calculation of:

$$Capacity = \frac{QAM_{BWTotal} - 2 * MAX_{BWSeSS}}{SS_{BW} * Concurrency}$$

Where *Capacity* is the number of customers a service group can handle defined by QAM bandwidth,  $QAM_{BWTotal}$  is the total usable QAM bandwidth available on the service group,  $MAX_{BWSeSS}$  is the max bandwidth used for a session (typically during session set up),  $SS_{BW}$  is the steady state bandwidth usage, and *Concurrency* is the concurrency rate of users in the service group.

This calculation leads to a capacity of approximately 720 customers per service group. There are several reasons why this is a low estimate.

1. The average usage is lower than an HD stream would dictate. This means more sessions can be allocated to a QAM than if they were all HD streams
2. The allocation of keeping max bandwidth in reserve per QAM is overkill. Not every session will use the max bandwidth on setup, so it is possible that less can be reserved for new sessions.
3. A 4% concurrency rate is on the high end of measured usages. Fluctuations to that value would be possible, but the majority of cases would be lower, and more users could be supported at lower concurrency rates.
4. The QRM uses smartmuxing which “Recycles” unused bandwidth in order to serve out more sessions than this analysis would indicate.

These factors provide significant leeway in capacity concerns. Operating at 720 customers per service group would be operationally ideal, but these factors would easily allow us to reach 800-1000 customers per service group without concerns to the customer experience. If we were to reach 1200-1500 customers per service group then the customer experience would begin to suffer due to Spec-Guide QAM bandwidth.

Customers per service group was measured through the CSM logs (but is also tracked in other ways by other groups) in order to determine if any service groups were over the relevant thresholds. Table 1 lists service groups to watch. It is noteworthy that the average number of customer per service group is INSERTNUMBERHERE, (which is BLAHBLAHBLAH lower than the first threshold) so overall we have plenty of QAM bandwidth to go around so long as it is properly allocated.

### 3 Conclusion

Stuff and things...

If sessions begin to overrun the Spec-Guide QAM bandwidth beyond what smartmuxing could handle, the customer experience would be affected in a few ways. First the QRM will send a few frames of lower quality for sessions in order to limit the bandwidth usage. Next, users would experience pixelation of the guide interface, delayed frames, and/or dropped frames in order to maintain the integrity of the Spec-Guide QAM. As a last resort the QRM will deny new sessions outright. The QRM will not allow sessions to be sent to a QAM once the bandwidth has been entirely allocated. All of these behaviours would be recorded in CSM logs.

Stuff and things...

There are no significant Spec-Guide QAM bandwidth related concerns.

The best way to ensure we are not overrunning our QAM bandwidth is to limit the number of customers per service group, and perform service group collapses and splits as needed.

Direct measurements of RF bandwidth in an automated fashion would be time consuming, costly, and unnecessary. The number of customers per service group serves the same purpose in this case, and requires no further monitoring.

## A Appendix A Example QRM Configuration

Stuff and things...

```
<resolution height="480">
<mpeg2>
<bitrateProfile name="Default_SD">1000000</bitrateProfile>
<bitrateProfile name="Low_SD">1000000</bitrateProfile>
<bitrateProfile name="MediumLow_SD">1200000</bitrateProfile>
<bitrateProfile name="Medium_SD">1500000</bitrateProfile>
<bitrateProfile name="High_SD">2000000</bitrateProfile>
<bitrate>4000000</bitrate>
<quant>3</quant>
<aquant>4</aquant>
<txtquant>2</txtquant>
</mpeg2>
<h264>
<bitrate>4000000</bitrate>
<quant>14</quant>
<aquant>14</aquant>
</h264>
```

```
<resolution height="720">
<mpeg2>
<bitrateProfile name="Default_HD">1600000</bitrateProfile>
<bitrateProfile name="Low_HD">1600000</bitrateProfile>
<bitrateProfile name="MediumLow_HD">1800000</bitrateProfile>
<bitrateProfile name="Medium_HD">2000000</bitrateProfile>
<bitrateProfile name="High_HD">2500000</bitrateProfile>
<bitrate>6000000</bitrate>
<quant>3</quant>
<aquant>5</aquant>
<txtquant>2</txtquant>
</mpeg2>
<h264>
<bitrate>10000000</bitrate>
<quant>10</quant>
<aquant>14</aquant>
</h264>
```

Market	SGID	Customer Count

Table 1: Service groups with customer counts over 720.

## B Appendix B Customers by SGID

CustPerSG/edprmnResults.txt

744 40019  
819 30003  
850 41010  
1006 40017  
1079 30006  
1083 30001

CustPerSG/knwdmiResults.txt

726 49011  
736 37025  
748 21032  
766 48076  
799 28022  
805 13010  
1222 48075

CustPerSG/mddcwiResults.txt

731 14034  
740 42028  
741 24011  
771 52031  
786 29002  
791 36021  
817 27017  
833 23004  
877 42035  
1014 13029  
1212 50002

CustPerSG/pldcorResults1.txt

729 40225  
739 40956  
755 14022  
763 40186  
766 20017  
779 20019  
789 40114  
794 40518  
802 15001  
806 27068  
815 27070  
828 14018  
834 27022  
836 40116  
842 20012  
843 20003  
908 14023  
991 14026  
992 40154  
1048 14030  
1049 14016  
1087 14035  
1413 14020  
1415 14019  
1878 14007

CustPerSG/pldcorResults2.txt

723 28011  
727 38023  
733 25005  
733 26003  
737 25002  
740 38034  
746 25006  
747 28009  
749 28008  
767 28006  
768 19002  
768 43017  
774 38035  
782 43016  
791 43013  
793 25019  
798 25020  
811 34004  
822 25008



827 19004  
837 43035  
843 43011  
849 25011  
858 25021  
864 19012  
880 26002  
882 25010  
901 34002  
902 43012  
921 25012  
940 38014  
959 19008  
1006 38030  
1014 25004  
1073 19007  
1077 34005  
1159 28002  
1200 43015  
1382 34003  
CustPerSG/renonvResults.txt  
721 11018  
728 11101  
730 12029  
730 14015  
741 12005  
741 12010  
749 10031  
749 11150  
752 17001  
757 16010  
760 10005  
760 11126  
766 14016  
769 11134  
770 16015  
775 17002  
782 11142  
789 11175  
797 11149  
797 11156  
797 12013  
804 11147  
804 15011  
813 11044  
815 12018

817 11141  
823 11068  
830 11094  
832 10050  
840 11151  
842 11045  
848 11165  
857 15022  
866 11114  
869 11174  
876 11179  
881 12020  
897 16001  
904 11140  
931 11123  
952 12024  
956 11059  
973 12025  
1008 17010  
1013 11144  
1017 12009  
1023 16016  
1038 11028  
1122 11104  
1208 12008  
CustPerSG/sldclaResults.txt - -  
CustPerSG/sldcmoResults.txt  
3628 24090  
CustPerSG/spdcscResults.txt  
720 61105  
734 26077  
737 22033  
737 26024  
752 24061  
757 26089  
758 32093  
782 30015  
794 22028