

Stitcher Capacity Analysis: Charter Internal Note

Brad Schoenrock
Video Operations Engineering
Charter Communications
Greenwood Village, CO

1 Introduction

Common wisdom within charter and Active Video was that stitchers with 96GB of ram had a capacity of 400 concurrent sessions per stitcher. That analysis was flawed. That conclusion was based on average session size, estimated concurrency rates, and the number of users on a market. This model leaves no room for variance, and in the event that a stitcher takes sessions of above average size (which will by definition happen half of the time) then that stitcher will go degraded. If calculations are based on concurrency instead of session size then variance in the concurrency must be taken in order to account for this effect, which originally was not the case.

Simulations of a stitcher in the lab were performed in the past which reinforced a confirmation bias. It is very likely that failed tests were dismissed for seemingly explainable reasons, and once success was obtained further testing was not conducted. In a tightly controlled lab environment it is possible to trick oneself into a false sense of security because 1) the size of sessions is more tightly controlled in the lab compared to in production, which is evidenced by the fact that the lab had to go out of their way to replicate the issue when prompted, and 2) the degradation of a stitcher is fundamentally a statistical event. With one stitcher in the lab it is possible to get lucky and not go degraded, even at session counts which are beyond estimated thresholds. In production, however, we have approximately 3350 stitchers which undergo peak prime time load every half hour for 5 hours, so we run that test 33500 times a day. If we actually find ourselves at that 400 session per stitcher that means we get 16750 degraded stitchers per day that have to be recovered. Clearly this is not operationally feasible. The question then becomes what are acceptable rates for stitcher degradation, and how many sessions can a stitcher take at that error rate.

A note on customer impact before we begin discussion of the analysis. Stitcher capacity being limited by memory will go through three phases. At first, markets approach capacity, and the result is that sessions undergo increased latency. As markets cross the capacity threshold and are underprovisioned users experience increased rates of guide unavailable. These effects are already being seen in production. When markets cross a higher threshold and

too many stitchers start going degraded all sessions then have to be redirected to other stitchers. When the remaining stitchers take that load they will be more likely go degraded under the increased load then they will go degraded themselves, and a cascading problem occurs. This will lead to a near 100% guide unavailable full market blackout.

2 Session Size Data Acquisition

Session size was measured with ansible and ps via the following command (example given for twcsc) -
ansible twcsc.spdc.sc-stitchers -m shell -a "ps -C html5client -o start, pid, etime, cmd, pcpu, rss, size" | tee -a twcsc-vca.txt
- which returns the status of all html5client processes currently running on that market. The results were a text file that must be parsed in order to extract session parameters. The parameters extracted were elapsed time running, CPU usage, RSS size, and SIZE as defined by the ps command. RSS and SIZE were converted to mb. Defunct responses from the ps command were discarded which account for < 5 sessions enterprise wide. The use of ansible means that not all stitchers respond without timeouts leading to not every stitcher responding. This was corrected for by comparing the number of successful ansible returns with the number of unreachable ansible returns. The ansible connection factor is 40%.

3 Session Size Analysis

Session parameters were loaded into python and the pandas utility was used in order to generate a statistical summary of sessions including mean, median, standard deviation, session counts, min, max, and quartiles. Visualization of that distribution was performed using matplotlib. Summary of session size for all markets aggregated and separated can be seen in appendix A.

4 Peak Usage Data Acquisition

Peak usage was determined through measurement of sessions reported in CSM logs. Each session's start, end, stitcher hit, and service group hit were collected and aggregated. Features of those log events were extracted and analysis has been performed to measure a multitude of CSM functionality including times of events, service group response times, session exit conditions, and more. Relevant to this analysis was session start times which have the feature for user login times and session lengths based on session start time and session exit time.

5 Peak Usage Measurement

User login times were collected for every CSM independantly, and each CSM represents about one third of the traffic on a market. Binned into 10 min intervals to ensure reasonable statistics and fitted prime time loads peak usage for each CSM was read from these histograms. In a 10 min interval a fraction of sessions are running at any given time, and based on measured session length from measured CSM data a prime time concurrent load can be calculated.

A distinction should be made between prime time load, and peak prime time load. A fit to the data will return prime time load which was used in this analysis, but every 30 minutes on the half hour a spike in usage occurs. Those spikes (peak prime time load) can be up to 2x the result from the fitted prime time load. If peak prime time analysis is preferred double the stitcher need.

6 Stitcher Capacity Analysis

blah blah blah ...

7 Market Capacity Analysis

blah blah blah ...

A Session Size Tables

All markets	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	7964	7964.000000	7964.000000	7964.00000
mean	0 days 16:57:29.043571	6.670191	704.325508	1016.00872
std	6 days 12:33:38.785822	17.375503	587.128283	751.04424
min	0 days 00:00:00	0.000000	26.296000	37.20400
25%	0 days 00:04:08	0.300000	306.063000	502.49400
50%	0 days 00:28:41.500000	0.900000	604.702000	906.51000
75%	0 days 01:13:26	4.500000	888.213000	1256.41700
max	131 days 03:35:10	145.000000	14712.656000	16320.25200

twctx	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	1234	1234.000000	1234.000000	1234.00000
mean	0 days 00:38:16.789303	4.593112	490.834882	751.81550
std	0 days 00:58:26.775946	10.083455	253.906020	352.73699
min	0 days 00:00:00	0.000000	25.388000	69.82800
25%	0 days 00:03:03.750000	0.300000	272.858000	446.80800
50%	0 days 00:18:51	0.900000	366.556000	588.50600
75%	0 days 00:53:25.750000	4.700000	670.100000	1007.64800
max	0 days 14:24:02	130.000000	1614.120000	2286.67600

spdcsc	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	1344	1344.000000	1344.000000	1344.00000
mean	0 days 05:33:22.727678	10.047545	747.331970	1148.42780
std	1 days 02:24:33.678990	22.922085	744.264176	1063.67384
min	0 days 00:00:00	0.000000	108.732000	163.03600
25%	0 days 00:03:07.500000	0.400000	363.873000	616.14000
50%	0 days 00:21:37.500000	1.300000	617.814000	876.55200
75%	0 days 01:02:09.250000	7.300000	818.191000	1253.15700
max	10 days 03:05:09	154.000000	11340.324000	12773.81600

mddcwi	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	663	663.000000	663.000000	663.00000
mean	0 days 01:13:31.853695	7.504827	871.828561	1323.55937
std	0 days 05:18:41.994081	17.288993	697.919157	1085.96406
min	0 days 00:00:01	0.000000	155.052000	222.77200
25%	0 days 00:03:22	0.400000	311.396000	528.56600
50%	0 days 00:21:19	1.300000	710.036000	1025.54400
75%	0 days 00:57:54.500000	6.350000	1140.840000	1598.61800
max	3 days 03:49:19	99.800000	5271.756000	8797.68000

bhnoh	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	38	38.000000	38.000000	38.00000
mean	45 days 04:09:45.289473	92.215789	578.135263	953.34589
std	41 days 20:07:34.367849	26.600442	98.453239	173.38133
min	0 days 00:07:49	1.200000	453.460000	751.18400
25%	28 days 09:15:25.250000	99.900000	542.662000	894.69800
50%	28 days 17:57:53	99.900000	556.336000	918.71400
75%	29 days 04:38:27.750000	99.900000	582.774000	981.54900
max	130 days 15:53:57	99.900000	1052.028000	1736.06000

twcsc	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	697	697.000000	697.000000	697.00000
mean	0 days 02:14:16.459110	5.776327	497.244746	757.31775
std	0 days 21:17:15.405321	13.323263	287.396410	387.90609
min	0 days 00:00:00	0.000000	36.252000	82.15600
25%	0 days 00:02:47	0.300000	273.080000	447.08000
50%	0 days 00:17:00	1.000000	347.376000	561.42400
75%	0 days 00:53:12	5.500000	693.360000	1035.06000
max	14 days 02:01:05	101.000000	2512.328000	3058.43600

slotca	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	189	189.000000	189.000000	189.00000
mean	2 days 05:08:17.068783	10.966667	517.918413	743.12943
std	12 days 06:00:47.279801	23.232777	254.689270	314.38865
min	0 days 00:00:00	0.000000	68.068000	113.12400
25%	0 days 00:02:09	0.400000	291.140000	466.89600
50%	0 days 00:13:11	1.900000	506.636000	705.66000
75%	0 days 00:58:16	8.800000	690.876000	947.54800
max	103 days 06:27:34	100.000000	1366.460000	1732.20800

ladcca	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	867	867.000000	867.000000	867.00000
mean	0 days 08:47:35.552479	8.360208	450.936166	669.28376
std	2 days 16:09:18.559904	18.904368	244.823500	307.58961
min	0 days 00:00:00	0.000000	12.792000	22.89200
25%	0 days 00:01:21.500000	0.400000	275.056000	466.51000
50%	0 days 00:09:43	1.400000	336.452000	557.32000
75%	0 days 00:47:16.500000	7.850000	594.040000	820.15200
max	28 days 05:24:52	150.000000	1996.924000	2973.56800

bhnfl	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	12	12.000000	12.000000	12.00000
mean	97 days 22:32:49	76.050000	488.248333	816.10533
std	59 days 00:42:48.484789	43.274063	128.220732	184.42619
min	0 days 00:02:23	0.200000	219.072000	375.20800
25%	97 days 20:51:46.250000	76.625000	451.802000	774.66400
50%	130 days 12:32:55	99.900000	489.430000	791.46800
75%	130 days 14:56:07.750000	100.000000	501.310000	862.90800
max	130 days 16:14:00	100.000000	797.216000	1138.86400

twcohd	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	611	611.000000	611.000000	611.00000
mean	0 days 00:41:26.962356	4.990344	499.298710	762.94798
std	0 days 00:55:17.385177	10.710850	253.252774	350.21950
min	0 days 00:00:00	0.000000	18.548000	55.67600
25%	0 days 00:03:01.500000	0.300000	282.944000	460.38400
50%	0 days 00:20:57	0.800000	373.420000	609.91200
75%	0 days 00:59:45	4.300000	683.328000	1013.02800
max	0 days 07:08:58	90.500000	1385.548000	2101.02800

sldcmo	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	811	811.000000	811.000000	811.00000
mean	0 days 00:48:12.890258	4.603083	991.308099	1350.01964
std	0 days 01:16:55.655732	9.059367	681.528005	798.25732
min	0 days 00:00:00	0.000000	140.636000	225.09600
25%	0 days 00:04:00	0.400000	402.364000	658.48400
50%	0 days 00:24:37	1.200000	858.612000	1207.46800
75%	0 days 00:58:13.500000	4.700000	1315.152000	1718.66600
max	0 days 16:32:18	93.000000	4124.472000	5313.78000

knwdmi	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	854	854.000000	854.000000	854.00000
mean	0 days 01:50:32.909836	7.179040	742.129691	1106.20582
std	0 days 11:45:10.901628	16.393025	730.700326	942.02897
min	0 days 00:00:01	0.000000	124.752000	209.78000
25%	0 days 00:02:39	0.400000	371.277000	626.71100
50%	0 days 00:18:28.500000	1.100000	638.800000	933.37400
75%	0 days 00:51:55	6.575000	920.667000	1273.83700
max	7 days 07:38:22	111.000000	13115.864000	14916.46000

bhnca	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	1	1.0	1.00	1.0
mean	0 days 00:00:06	30.5	247.74	404.7
std	NaN	NaN	NaN	NaN
min	0 days 00:00:06	30.5	247.74	404.7
25%	0 days 00:00:06	30.5	247.74	404.7
50%	0 days 00:00:06	30.5	247.74	404.7
75%	0 days 00:00:06	30.5	247.74	404.7
max	0 days 00:00:06	30.5	247.74	404.7

twcnysydc	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	442	442.000000	442.000000	442.00000
mean	0 days 01:01:03.726244	6.148190	500.684950	763.87129
std	0 days 05:46:28.356034	15.162467	260.041742	358.83244
min	0 days 00:00:00	0.000000	84.572000	132.11200
25%	0 days 00:03:06	0.300000	276.576000	455.65000
50%	0 days 00:22:50	0.800000	367.764000	581.14200
75%	0 days 00:59:00.750000	4.875000	697.221000	1039.71500
max	3 days 13:49:05	118.000000	1466.484000	1944.98800

sldcla	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	358	358.000000	358.000000	358.00000
mean	0 days 00:58:12.290502	6.725419	680.787911	928.31454
std	0 days 05:02:56.766229	14.714847	384.019453	454.63616
min	0 days 00:00:00	0.000000	24.804000	55.67600
25%	0 days 00:01:33	0.300000	314.110000	508.44400
50%	0 days 00:16:41	1.300000	614.728000	848.91000
75%	0 days 00:54:09.250000	6.300000	965.019000	1233.39500
max	3 days 22:07:52	105.000000	1957.948000	2583.84800

edprmn	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	703	703.000000	703.000000	703.00000
mean	3 days 23:41:03.660028	9.510953	687.757826	1037.11116
std	14 days 10:19:40.838831	19.807971	875.313315	1095.16143
min	0 days 00:00:00	0.000000	26.296000	37.20400
25%	0 days 00:01:28	0.400000	299.240000	519.00800
50%	0 days 00:13:56	1.700000	563.340000	828.87200
75%	0 days 00:52:57	9.950000	807.120000	1152.59800
max	61 days 20:24:45	136.000000	12119.260000	13713.00400

bhnal	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	703	703.000000	703.000000	703.00000
mean	3 days 23:41:03.660028	9.510953	687.757826	1037.11116
std	14 days 10:19:40.838831	19.807971	875.313315	1095.16143
min	0 days 00:00:00	0.000000	26.296000	37.20400
25%	0 days 00:01:28	0.400000	299.240000	519.00800
50%	0 days 00:13:56	1.700000	563.340000	828.87200
75%	0 days 00:52:57	9.950000	807.120000	1152.59800
max	61 days 20:24:45	136.000000	12119.260000	13713.00400

twcnny.nyd	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	439	439.000000	439.000000	439.00000
mean	0 days 00:42:33.881548	3.782232	505.566378	770.06229
std	0 days 00:56:37.780297	10.605681	253.195549	353.68918
min	0 days 00:00:00	0.000000	84.100000	132.11200
25%	0 days 00:05:23.500000	0.200000	284.092000	456.28800
50%	0 days 00:25:12	0.700000	413.500000	646.48800
75%	0 days 00:54:22.500000	2.350000	682.940000	1027.04000
max	0 days 07:27:59	119.000000	1451.068000	1988.59200

renonv	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	990	990.000000	990.000000	990.00000
mean	0 days 01:00:09.429292	4.826869	968.266428	1315.58436
std	0 days 01:28:07.999288	11.318484	643.358079	761.49195
min	0 days 00:00:00	0.000000	87.840000	156.11200
25%	0 days 00:04:56.250000	0.400000	407.672000	640.68700
50%	0 days 00:32:28.500000	0.900000	857.420000	1223.10600
75%	0 days 01:13:59.750000	4.000000	1267.769000	1644.98900
max	0 days 17:49:29	100.000000	4509.704000	5523.24800

dldctx	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	1240	1240.000000	1240.000000	1240.00000
mean	0 days 09:59:56.741129	9.365484	484.178539	729.23009
std	3 days 22:37:16.728882	18.061337	295.350663	359.44808
min	0 days 00:00:00	0.000000	84.692000	132.11200
25%	0 days 00:01:12	0.600000	274.080000	477.95400
50%	0 days 00:04:50.500000	3.300000	343.756000	568.51400
75%	0 days 00:38:54	9.250000	635.297000	909.19700
max	70 days 04:14:55	119.000000	2011.240000	2581.69200

bhdcal	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	449	449.000000	449.000000	449.00000
mean	3 days 10:05:22.552338	12.213586	511.082013	745.66570
std	15 days 10:04:40.019196	27.009075	278.026679	344.30193
min	0 days 00:00:00	0.000000	107.364000	158.71200
25%	0 days 00:03:04	0.500000	296.092000	492.30400
50%	0 days 00:22:34	1.400000	424.240000	654.30800
75%	0 days 01:02:13	7.300000	660.296000	913.29200
max	116 days 16:49:08	115.000000	1892.232000	2884.86400

twchi	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	0	0	0	0
unique	0	0	0	0

pldcor	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	389	389.000000	389.000000	389.00000
mean	1 days 05:24:52.033419	11.805141	815.620380	1175.33370
std	5 days 16:38:37.442890	24.115209	733.513476	1041.47454
min	0 days 00:00:01	0.000000	125.612000	211.30800
25%	0 days 00:01:08	0.300000	330.576000	576.40000
50%	0 days 00:14:33	2.300000	636.780000	919.36800
75%	0 days 01:38:37	8.500000	1054.984000	1405.63200
max	37 days 14:45:31	99.900000	6727.340000	8748.13200

bodcma	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	593	593.000000	593.000000	593.00000
mean	0 days 00:33:41.399662	5.413322	537.107386	765.38517
std	0 days 00:44:57.083644	12.469223	307.470807	369.60064
min	0 days 00:00:00	0.000000	144.824000	210.89600
25%	0 days 00:01:55	0.300000	280.832000	462.84000
50%	0 days 00:17:30	1.000000	412.584000	637.72800
75%	0 days 00:49:06	5.200000	722.672000	989.54400
max	0 days 08:33:57	153.000000	2237.936000	2545.99200

twcca	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	505	505.000000	505.000000	505.00000
mean	0 days 00:35:02.885148	4.480000	515.805513	793.22996
std	0 days 00:45:24.634973	8.867115	255.678862	365.21607
min	0 days 00:00:00	0.000000	122.608000	176.76000
25%	0 days 00:03:25	0.300000	288.092000	459.67200
50%	0 days 00:18:04	0.900000	398.276000	666.70800
75%	0 days 00:50:38	4.200000	702.064000	1059.37600
max	0 days 05:29:07	82.500000	1365.628000	1919.64400

nvdctn	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	773	773.000000	773.000000	773.00000
mean	1 days 23:34:24.972833	9.574774	581.608310	834.08841
std	11 days 21:42:02.836950	21.275061	351.990133	428.15494
min	0 days 00:00:00	0.000000	55.104000	105.19200
25%	0 days 00:02:02	0.400000	299.696000	492.51200
50%	0 days 00:19:12	1.500000	499.720000	732.01600
75%	0 days 00:56:50	7.900000	779.652000	1067.57600
max	127 days 05:50:07	115.000000	2500.564000	3165.04800

blngmt	ELAPSED	CPU%	RSS(mb)	SIZE(mb)
count	269	269.000000	269.000000	269.00000
mean	0 days 00:37:08.981412	5.629368	547.594528	775.47467
std	0 days 00:54:32.902980	9.371941	439.892380	541.74936
min	0 days 00:00:00	0.000000	122.984000	177.21200
25%	0 days 00:01:03	0.400000	275.140000	429.76800
50%	0 days 00:14:04	1.300000	330.268000	515.21600
75%	0 days 00:54:23	7.800000	646.032000	951.88400
max	0 days 05:50:57	67.500000	2521.860000	3125.04800

B LaTeX example syntax IGNORE

1
hey
2
there
3
listen

1. First thing
2. Next thing
3. One last thing

$\alpha\beta\gamma\delta\epsilon\zeta\eta\theta\iota\kappa\lambda\mu\nu\xi\omicron\pi\rho\sigma\tau\upsilon\phi\chi\psi\omega$
weil@math.msu.edu
This is a short example of a x^2 document with

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

an out of paragraph equation.
this is an in text citation. [1]

References

- [1] B.Schoenrock, me., **1** (2018)